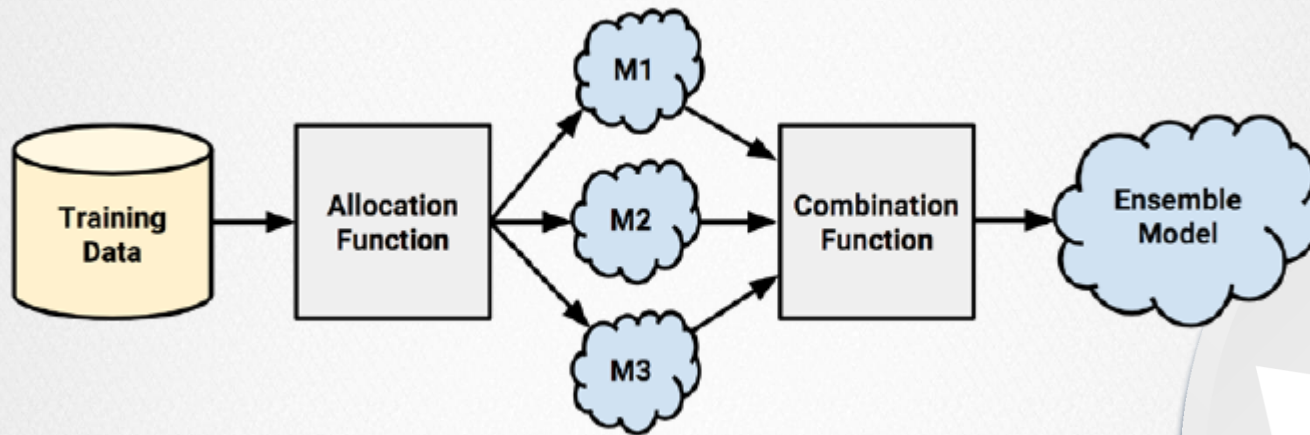


# Model Ensembles

**Practical Machine Learning (with R)**  
UC Berkeley

# Ensembles



Source: *Machine Learning with R*

# TWO BIG IDEAS

## ➤ **Wisdom of the crowds**

It is better to make estimates from multiple models (**ensembles**) than individual models

- Better predictions
- Lower variance for the same model

## ➤ **Greed is bad. Patience is good.**

It is better to slowly approach your solution than arrive at an answer directly.

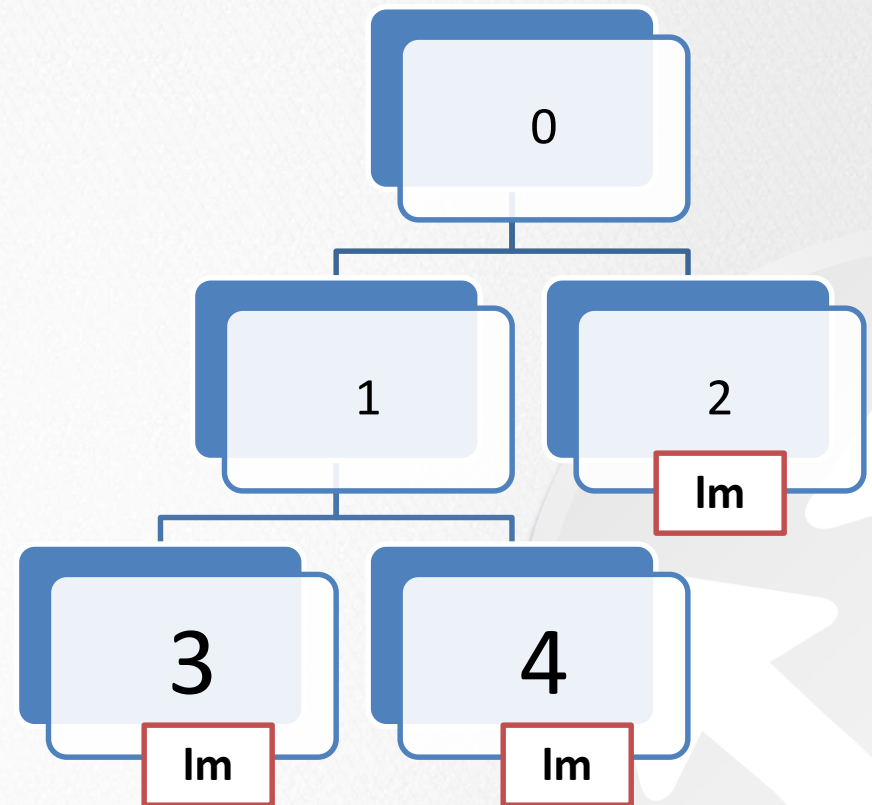
- More accurate solutions





# Tree Enhancement: M5

- **Wisdom of the Crowd!**
- Having one value represent the entirety of the node leaves information in the node.
- Function in the node is a simple average
- Use something better
  - **M5** put linear models in nodes of trees

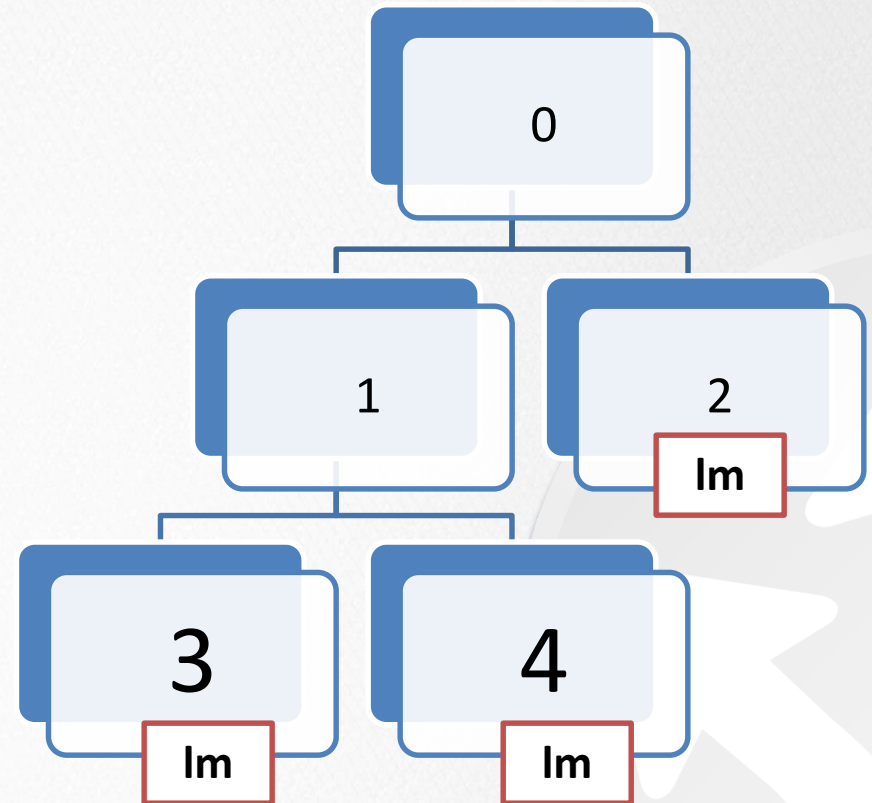



# M5 Tree Enhancement (cont.)

## → Greed is bad

- linear models are built on the residuals of the tree model.

- Models are recursive




# BAGGING MODELS

➔ Brieman:

"Bagging is a general approach that uses bootstrapping in conjunction with any regression (or classification) model to construct an ensemble."

```
1 for  $i = 1$  to  $m$  do
2   |   Generate a bootstrap sample of the original data
3   |   Train an unpruned tree model on this sample
4 end
```

$$\hat{y} = \frac{\sum_i \hat{y}_i}{m}$$





# BAGGING NOTES

## ➤ Lowers variance

- Increases stability
- Has less effect on lower variance models (e.g. linear models)
- More effect on weak learners

## ➤ Disadvantages

- Computational cost → but parallelizable
- Reduces Interpretability





# RANDOM FOREST

- **Wisdom of the Crowds:** Bagging
- **Greed is bad:** consider subset of predictors at each split

```
1 Select the number of models to build,  $m$ 
2 for  $i = 1$  to  $m$  do
3     Generate a bootstrap sample of the original data
4     Train a tree model on this sample
5     for each split do
6         Randomly select  $k$  ( $< P$ ) of the original predictors
7         Select the best predictor among the  $k$  predictors and
           partition the data
8     end
9     Use typical tree model stopping criteria to determine when a
       tree is complete (but do not prune)
10 end
```

# TUNING PARAMETER

$m_{\text{try}}$  : number of predictors to use at each split

- **regression** 1/3rd of number predictors
  - **classification**  $\sqrt{\text{number of predictors}}$
- ➔ Kuhn: “Starting with five values of  $k$  that are somewhat evenly spaced across the range from 2 to  $P$ ”.



# Random Forests

| Strengths   | Weaknesses  |
|---|---|
| <ul style="list-style-type: none"><li>• An all-purpose model that performs well on most problems</li><li>• Can handle noisy or missing data as well as categorical or continuous features</li><li>• Selects only the most important features</li><li>• Can be used on data with an extremely large number of features or examples</li></ul> | <ul style="list-style-type: none"><li>• Unlike a decision tree, the model is not easily interpretable</li><li>• May require some work to tune the model to the data</li></ul> |



# ADVANTAGES

- ➔ No overfitting
- ➔ More trees better (limited by computation time/power only)
- ➔ In caret, parameters are considered independently
- ➔ Because each learner is selected independently of all previous learners, Random Forests is robust to a noisy response
- ➔ Computationally efficient -- each tree built on subset of predictors at each split.
- ➔ Use any tree variants as "base learner": CART, ctree, etc



# BOOSTING



# BOOSTING

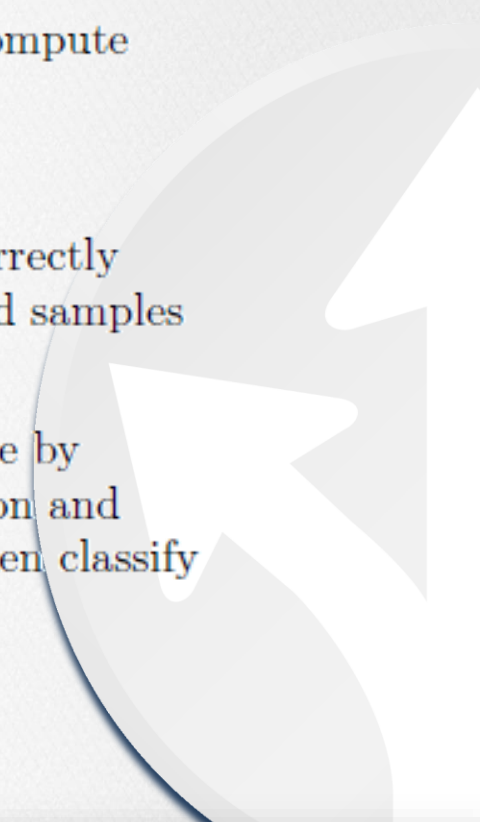
- ⇒ Single models work;
  - Multiple models work better
- ⇒ Idea is simple:
  - **Fit first** model:  $\hat{y}_1 \sim f_1(x)$
  - **Fit** errors/residuals: 
$$\begin{aligned}\hat{y}_2 &= f_2(y - \hat{y}_1) \\ &= f_2(y - f_1(x)) \\ &= f_2(x)\end{aligned}$$
  - **Iterate:**  $\hat{y}_i = (y - \hat{y}_{i-1}) \sim f_i(x)$
  - **Predict:**  $\hat{y} \sim \sum_i f_i(x)$






# ADABOOST (SHAPIORE/FREUND)

- 1 Let one class be represented with a value of  $+1$  and the other with a value of  $-1$
- 2 Let each sample have the same starting weight ( $1/n$ )
- 3 **for**  $k = 1$  **to**  $K$  **do**
- 4     Fit a weak classifier using the weighted samples and compute the  $k$ th model's misclassification error ( $err_k$ )
- 5     Compute the  $k$ th stage value as  $\ln((1 - err_k) / err_k)$ .
- 6     Update the sample weights giving more weight to incorrectly predicted samples and less weight to correctly predicted samples
- 7 **end**
- 8 Compute the boosted classifier's prediction for each sample by multiplying the  $k$ th stage value by the  $k$ th model prediction and adding these quantities across  $k$ . If this sum is positive, then classify the sample in the  $+1$  class, otherwise the  $-1$  class.



# Adaboost Stage Value



# BOOSTING NOTES

- ➔ Additive models
- ➔ Works best with “weak learners”
  - i.e. ungreedy, low bias, low variance
  - ~~Any~~ Most models with a tuning parameter can be a weak learner
  - Trees are excellent weak learners
    - Weak → “restricted depth”
- ➔ Residuals or errors define a gradient
- ➔ Interpreted as forward step-wise regression with exponential loss





# SIMPLE GRADIENT BOOSTING (REGRESSION)

- 1 Select tree depth,  $D$ , and number of iterations,  $K$
- 2 Compute the average response,  $\bar{y}$ , and use this as the initial predicted value for each sample
- 3 **for**  $k = 1$  **to**  $K$  **do**
  - 4 Compute the residual, the difference between the observed value and the *current* predicted value, for each sample
  - 5 Fit a regression tree of depth,  $D$ , using the residuals as the response
  - 6 Predict each sample using the regression tree fit in the previous step
  - 7 Update the predicted value of each sample by adding the previous iteration's predicted value to the predicted value generated in the previous step
- 8 **end**

Naïve  
Guess

# SIMPLE GRADIENT BOOSTING (CLASSIFICATION)

- 1 Initialized all predictions to the sample log-odds:  $f_i^{(0)} = \log \frac{\hat{p}}{1-\hat{p}}$ .
- 2 **for** *iteration*  $j = 1 \dots M$  **do**
- 3     Compute the residual (i.e. gradient)  $z_i = y_i - \hat{p}_i$
- 4     Randomly sample the training data
- 5     Train a tree model on the random subset using the residuals as the outcome
- 6     Compute the terminal node estimates of the Pearson residuals:  
$$r_i = \frac{1/n \sum_i^n (y_i - \hat{p}_i)}{1/n \sum_i^n \hat{p}_i (1 - \hat{p}_i)}$$
- 7     Update the current model using  $f_i = f_i + \lambda f_i^{(j)}$
- 8 **end**





# STOCHASTIC GRADIENT BOOSTING

## ⇒ Gradient Boosting Susceptible to Overfitting

- Apply “regularization/shrinkage”
  - Use  $\lambda$  (“Learning Rate”)  
Rather than add the entirety of the residuals, add a fraction of the residuals at each iteration.

$$\hat{y} \sim \lambda \sum_i f_i(x) \quad 0 < \lambda \leq 1$$

- Small values for  $\lambda$  ( $\sim 0.01$ ) work best
- $\lambda \sim 1/\text{computational time} \sim 1/\text{storage size}$

## ⇒ Use bagging, as well

- Bagging Fraction: a sample of data in each loop iteration





# APPENDIX

